

RabbitMQ / Bunny

RubyFTW

Tuesday, September 23, 2014

Dan Swain @dantswain <http://dantswain.com> dan.t.swain@gmail.com

Real Time Bidding Lead
simpli.fi

About Simpli.fi

- Simpli.fi is a real-time bidding (RTB) platform
 - ~500,000 queries per second
 - < 50 msec latency
 - C++ bidder, Ruby infrastructure, Erlang services
- We use RabbitMQ at Simpli.fi
 - Relatively new to us
 - Distribute thousands of messages per second across multiple data centers



Why use RabbitMQ?

- Message-based architecture is magical
 - Not the same way that Rails is “magical”
- Asynchronous request processing with built-in routing and queueing
- Trivially share data across multiple applications (pub/sub)
 - Potentially written in multiple languages
- Trivially scale up the number of workers
- Fault tolerance!
- No polling

Outline


- RabbitMQ concepts
- Ruby Bunny
- Example



-
- Advanced Message Queue Protocol
 - JP Morgan Chase, 2003
 - Implementations of AMQP 1.0
 - SwiftMQ - Java (JMS)
 - Windows Azure Service Bus
 - Apache - Qpid, ActiveMQ, Apollo
 - IBM MQ Light
 - **RabbitMQ**



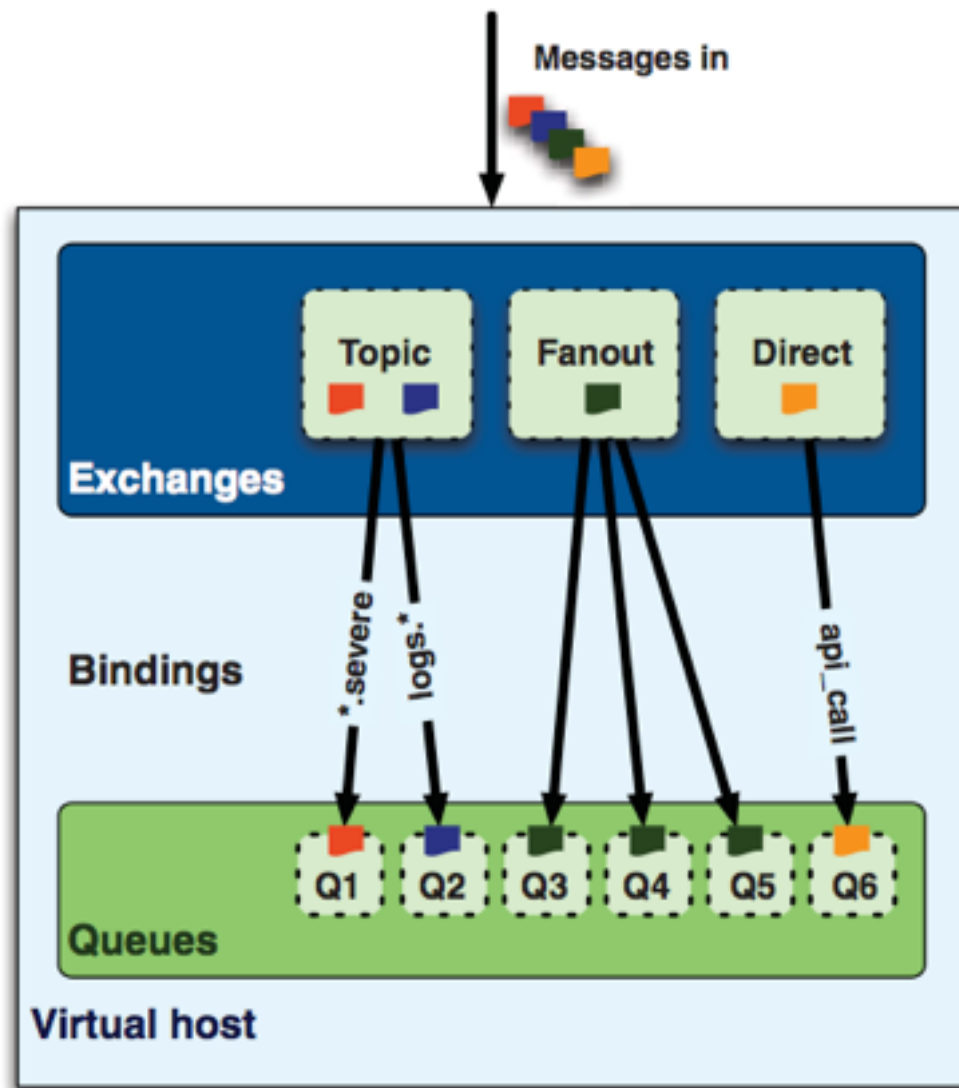
- Rabbit Technologies, 2007 -> VMWare, 2010 -> Pivotal, 2013
- Open Source, written in Erlang
- Current - 3.3.5
- Also supports ZeroMQ
- Scalable by clustering*
- Failover support
- RabbitMQ In Action - Videla, Williams (Manning)





- brew install rabbitmq / apt-get install rabbitmq-server
- rabbitmq-server -detached / rabbitmqctl stop
- Management Plugin - HTTP / REST - VERY useful -
 - rabbitmq-plugins enable rabbitmq_management
 - http://localhost:15672 - guest/guest
- rabbitmqadmin - Python script controls via REST
- Default username/password is guest/guest
 - Only works via 127.0.0.1





- Exchanges
 - Publish to exchanges
 - Routing logic
- Queues
 - Subscribe to queues
 - **Bind** queues to Exchanges
- Exchange Types: Direct, Fanout, Topic
- Queues and messages can be durable
- Virtual hosts - separate concerns

amqp uri

amqp://user:password@host/vhost

amqp://user:password@host <- default vhost = "/"

Distributing the rabbits

- Clustering intended for same-data center usage
- Two options: Shovels and Federation
- Shovels
 - More manual
 - More flexible
- Federation
 - More automated
 - Exchange has the appearance of existing on each broker
- github.com:dantswain/rabbitmq-federation-example

Rubying the rabbits: Bunny

- `gem install bunny`
- <http://rubybunny.info> - <https://github.com/ruby-amqp/bunny>
- rubybunny.info - “Guides”
- “Real” API doc site: <http://reference.rubybunny.info/>
- Official RabbitMQ tutorials using Bunny:
<http://www.rabbitmq.com/tutorials/tutorial-one-ruby.html>

Design patterns

Pedagogical

1 "Hello World!"

The simplest thing that does *something*



Python | Java | Ruby | PHP
| C#

(Use Topics)

2 Work queues

Distributing tasks among workers

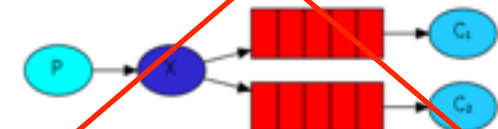


Python | Java | Ruby | PHP
| C#

Use Topics

3 Publish/Subscribe

Sending messages to many consumers at once



Python | Java | Ruby | PHP
| C#

4 Routing

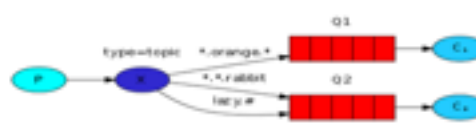
Receiving messages selectively



Python | Java | Ruby | PHP
| C#

5 Topics

Receiving messages based on a pattern



Python | Java | Ruby | PHP
| C#

6 RPC

Remote procedure call implementation



Python | Java | Ruby | PHP
| C#

Use Topics

Good luck...

Lessons so far

- Everything looks like topics. It's OK.
- Let consumers ensure the queue/binding exist.
- Federation was tricky to figure out, but works great.
 - Caveat: If a node goes down, prepare for big queues.
- “Redeliver” doesn't necessarily mean what you think it means. Re-publishing always works.
- When queues get big, RabbitMQ gets slower, and can take the box with it.
- Give yourself an easy way to spawn extra workers.
- Use monit (or god, bluepill, whatever) to make sure workers come back up without intervention, but check that they are ACTUALLY coming back up (Groundhog Day suicide effect).